

LangChain Expression Language

Streaming

Streaming With LangChain

Streaming is critical in making applications based on LLMs feel responsive to end-users.

Important LangChain primitives like LLMs, parsers, prompts, retrievers, and agents implement the LangChain Runnable Interface.

This interface provides two general approaches to stream content:

- sync stream and async astream: a default
 implementation of streaming that streams the final
 output from the chain.
- 2. async astream_events and async astream_log: these provide a way to stream both intermediate steps and final output from the chain.

Using Stream

All Runnable objects implement a sync method called stream and an async variant called astream.

These methods are designed to stream the final output in chunks, yielding each chunk as soon as it is available.

Streaming is only possible if all steps in the program know how to process an **input stream**; i.e., process an input chunk one at a time, and yield a corresponding output chunk.

The complexity of this processing can vary, from straightforward tasks like emitting tokens produced by an LLM, to more challenging ones like streaming parts of JSON results before the entire JSON is complete.

The best place to start exploring streaming is with the single most important components in LLMs apps— the LLMs themselves!

LLMs and Chat Models

Large language models and their chat variants are the primary bottleneck in LLM based apps. (2)

Large language models can take **several seconds** to generate a complete response to a query. This is far slower than the **~200-300 ms** threshold at which an application feels responsive to an end user.

The key strategy to make the application feel more responsive is to show intermediate progress; e.g., to stream the output from the model **token by token**.

```
from langchain.chat_models import
ChatAnthropic

model = ChatAnthropic()

chunks = []
async for chunk in model.astream("hello. tell
me something about yourself"):
    chunks.append(chunk)
    print(chunk.content, end="|", flush=True)
```

```
Hello|!| My| name| is| Claude|.| I|'m| an|
AI| assistant| created| by| An|throp|ic| to|
be| helpful|,| harmless|,| and| honest|.||
```

Let's inspect one of the chunks

```
chunks[0]
```

AIMessageChunk(content=' Hello')

We got back something called an AIMessageChunk. This chunk represents a part of an AIMessage.

Message chunks are additive by design – one can simply add them up to get the state of the response so far!

```
chunks[0] + chunks[1] + chunks[2] + chunks[3]
+ chunks[4]
```

AIMessageChunk(content=' Hello! My name is')

Chains

Virtually all LLM applications involve more steps than just a call to a language model.

Let's build a simple chain using LangChain Expression Language (LCEL) that combines a prompt, model and a parser and verify that streaming works.

We will use StrOutputParser to parse the output from the model. This is a simple parser that extracts the content field from an AIMessageChunk, giving us the token returned by the model.



LCEL is a *declarative* way to specify a "program" by chaining together different LangChain primitives.

Chains created using LCEL benefit from an automatic implementation of stream, and astream allowing streaming of the final output. In fact, chains created with LCEL implement the entire standard Runnable interface.

```
from langchain_core.output_parsers import
StrOutputParser
from langchain_core.prompts import
ChatPromptTemplate

prompt =
ChatPromptTemplate.from_template("tell me a
  joke about {topic}")
parser = StrOutputParser()
chain = prompt | model | parser

async for chunk in chain.astream({"topic":
  "parrot"}):
    print(chunk, end="|", flush=True)
```

Sure|,| here|'s| a| funny| joke| about| a| par|rot|:|

Why| doesn|'t| a| par|rot| ever| get| hungry|
at| night|?| Because| it| has| a| light|
snack| before| bed|!||

(i) NOTE

You do not have to use the LangChain Expression
Language to use LangChain and can instead rely on a
standard imperative programming approach by caling
invoke, batch or stream on each component
individually, assigning the results to variables and then
using them downstream as you see fit.

If that works for your needs, then that's fine by us 👌!

Working with Input Streams

What if you wanted to stream JSON from the output as it was being generated?

If you were to rely on json. loads to parse the partial json, the parsing would fail as the partial json wouldn't be valid json.

You'd likely be at a complete loss of what to do and claim that it wasn't possible to stream JSON.

Well, turns out there is a way to do it – the parser needs to operate on the **input stream**, and attempt to "auto-complete" the partial json into a valid state.

Let's see such a parser in action to understand what this means.

```
from langchain_core.output_parsers import
JsonOutputParser
from langchain_openai.chat_models import
ChatOpenAI
model = ChatOpenAI()
chain = model | JsonOutputParser() # This
parser only works with OpenAI right now
async for text in chain.astream(
    'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries" which contains a list of
countries. Each country should have the key
`name` and `population`'
):
    print(text, flush=True)
```

```
{}
{'countries': []}
```

```
{'countries': [{}]}
{'countries': [{'name': ''}]}
{'countries': [{'name': 'France'}]}
{'countries': [{'name': 'France',
'population': ''}]}
{'countries': [{'name': 'France',
'population': '67'}]}
{'countries': [{'name': 'France',
'population': '67,'}]}
{'countries': [{'name': 'France',
'population': '67,022'}]}
{'countries': [{'name': 'France',
'population': '67,022,'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name': ''}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': ''}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,'}]}
```

```
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'}, {}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'},
{'name': ''}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'},
{'name': 'Japan'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'},
{'name': 'Japan', 'population': ''}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'},
{'name': 'Japan', 'population': '126'}]}
{'countries': [{'name': 'France',
'population': '67,022,000'}, {'name':
'Spain', 'population': '46,754,784'},
```

```
{'name': 'Japan', 'population': '126,'}]}
{'countries': [{'name': 'France',
    'population': '67,022,000'}, {'name':
    'Spain', 'population': '46,754,784'},
    {'name': 'Japan', 'population': '126,860'}]}
{'countries': [{'name': 'France',
    'population': '67,022,000'}, {'name':
    'Spain', 'population': '46,754,784'},
    {'name': 'Japan', 'population': '126,860,'}]}
{'countries': [{'name': 'France',
    'population': '67,022,000'}, {'name':
    'Spain', 'population': '46,754,784'},
    {'name': 'Japan', 'population':
    '126,860,301'}]}
```

Now, let's **break** streaming. We'll use the previous example and append an extraction function at the end that extracts the country names from the finalized JSON.



DANGER

Any steps in the chain that operate on **finalized inputs** rather than on **input streams** can break streaming functionality via stream or astream.



TIP

Later, we will discuss the astream_events API which streams results from intermediate steps. This API will stream results from intermediate steps even if the chain contains steps that only operate on **finalized inputs**.

```
from langchain_core.output_parsers import
JsonOutputParser
# A function that operates on finalized
inputs
# rather than on an input_stream
def _extract_country_names(inputs):
    """A function that does not operates on
input streams and breaks streaming."""
    if not isinstance(inputs, dict):
        return ""
    if "countries" not in inputs:
        return ""
    countries = inputs["countries"]
    if not isinstance(countries, list):
        return ""
    country_names = [
        country.get("name") for country in
countries if isinstance(country, dict)
```

```
return country_names

chain = model | JsonOutputParser() |
_extract_country_names

async for text in chain.astream(
    'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries"'
):
    print(text, end="|", flush=True)
```

[None, None, None]|

Generator Functions

Le'ts fix the streaming using a generator function that can operate on the **input stream**.



A generator function (a function that uses yield) allows writing code that operators on **input streams**

from langchain_core.output_parsers import JsonOutputParser async def _extract_country_names_streaming(input_stream): """A function that operates on input streams.""" country_names_so_far = set() async for input in input_stream: if not isinstance(input, dict): continue if "countries" not in input: continue countries = input["countries"] if not isinstance(countries, list): continue for country in countries: name = country.get("name") if not name: continue if name not in country_names_so_far: yield name

country_names_so_far.add(name)

```
chain = model | JsonOutputParser() |
_extract_country_names_streaming

async for text in chain.astream(
    'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries"'
):
    print(text, end="|", flush=True)
```

France|Spain|Japan|

Non-streaming components

Some built-in components like Retrievers do not offer any streaming. What happens if we try to stream them? ••

```
from langchain_community.vectorstores import
FAISS
from langchain_core.output_parsers import
StrOutputParser
from langchain_core.prompts import
ChatPromptTemplate
from langchain_core.runnables import
RunnablePassthrough
```

```
from langchain_openai import OpenAIEmbeddings
template = """Answer the question based only
on the following context:
{context}
Question: {question}
11 11 11
prompt =
ChatPromptTemplate.from_template(template)
vectorstore = FAISS.from_texts(
    ["harrison worked at kensho", "harrison
likes spicy food"],
    embedding=OpenAIEmbeddings(),
)
retriever = vectorstore.as_retriever()
chunks = [chunk for chunk in
retriever.stream("where did harrison work?")]
chunks
```

```
[[Document(page_content='harrison worked at
kensho'),
   Document(page_content='harrison likes spicy
food')]]
```

Stream just yielded the final result from that component.



An LCEL chain constructed using using non-streaming components, will still be able to stream in a lot of cases, with streaming of partial output starting after the last non-streaming step in the chain.

```
for chunk in retrieval_chain.stream(
    "Where did harrison work? " "Write 3 made
up sentences about this place."
```

```
04/02/2024, 19:51
```

):

```
print(chunk, end="|", flush=True)
```

Streaming | 🦜 🔗 Langchain

```
|H|arrison| worked| at| Kens|ho|,| a|
renowned| technology| company| known| for|
revolution|izing| the| artificial|
intelligence| industry|.
|K|ens|ho|,| located| in| the| heart| of|
Silicon| Valley|,| is| famous| for| its|
cutting|-edge| research| and| development|
in| machine| learning|.
|With| its| state|-of|-the|-art| facilities|
and| talented| team|,| Kens|ho| has| become|
a| hub| for| innovation| and| a| sought|-
after| workplace| for| tech| enthusiasts|
like| Harrison|.||
```

Now that we've seen how stream and astream work, let's venture into the world of streaming events.

Using Stream Events

Event Streaming is a **beta** API. This API may change a bit based on feedback.



Introduced in langchain-core 0.1.14.

```
import langchain_core
langchain_core.__version__
'0.1.14'
```

For the astream_events API to work properly:

- Use async throughout the code to the extent possible (e.g., async tools etc)
- Propagate callbacks if defining custom functions / runnables
- Whenever using runnables without LCEL, make sure to call .astream() on LLMs rather than .ainvoke to force the LLM to stream tokens.
- Let us know if anything doesn't work as expected! :)

Event Reference

Below is a reference table that shows some events that might be emitted by the various Runnable objects.

(i) NOTE

When streaming is implemented properly, the inputs to a runnable will not be known until after the input stream has been entirely consumed. This means that inputs will often be included only for end events and rather than for start events.

event	name	
on_chat_model_start	[model name]	
on_chat_model_stream	[model name]	AlMessageCh
on_chat_model_end	[model name]	
on_llm_start	[model name]	
on_llm_stream	[model name]	'Hello'
on_llm_end	[model name]	

event	name	
on_chain_start	format_docs	
on_chain_stream	format_docs	"hello world!
on_chain_end	format_docs	
on_tool_start	some_tool	
on_tool_stream	some_tool	{"x":1, "y": "2
on_tool_end	some_tool	
on_retriever_start	[retriever name]	
on_retriever_chunk	[retriever name]	{documents:
on_retriever_end	[retriever name]	
on_prompt_start	[template_name]	
on_prompt_end	[template_name]	

Chat Model

Let's start off by looking at the events produced by a chat model.

```
events = []
async for event in
model.astream_events("hello", version="v1"):
    events.append(event)
```

/home/eugene/.pyenv/versions/3.11.4/envs/langch packages/langchain_core/_api/beta_decorator.py: API is in beta and may change in the future. warn_beta(

(i) NOTE

Hey what's that funny version="v1" parameter in the API?!



This is a **beta API**, and we're almost certainly going to make some changes to it.

This version parameter will allow us to mimimize such breaking changes to your code.

In short, we are annoying you now, so we don't have to annoy you later.

Let's take a look at the few of the start event and a few of the end events.

```
events[:3]
```

```
[{'event': 'on_chat_model_start',
  'run id': 'd78b4ffb-0eb1-499c-8a90-
8e4a4aa2edae',
  'name': 'ChatOpenAI',
  'tags': [],
  'metadata': {},
  'data': {'input': 'hello'}},
 {'event': 'on_chat_model_stream',
  'run_id': 'd78b4ffb-0eb1-499c-8a90-
8e4a4aa2edae',
  'tags': [],
  'metadata': {},
  'name': 'ChatOpenAI',
  'data': {'chunk':
AIMessageChunk(content='')}},
 {'event': 'on_chat_model_stream',
  'run id': 'd78b4ffb-0eb1-499c-8a90-
8e4a4aa2edae',
  'tags': [],
  'metadata': {},
```

```
'name': 'ChatOpenAI',
  'data': {'chunk':
AIMessageChunk(content='Hello')}}]
```

```
events[-2:]
```

```
[{'event': 'on_chat_model_stream',
  'run id': 'd78b4ffb-0eb1-499c-8a90-
8e4a4aa2edae',
  'tags': [],
  'metadata': {},
  'name': 'ChatOpenAI',
  'data': {'chunk':
AIMessageChunk(content='')}},
 {'event': 'on_chat_model_end',
  'name': 'ChatOpenAI',
  'run id': 'd78b4ffb-0eb1-499c-8a90-
8e4a4aa2edae',
  'tags': [],
  'metadata': {},
  'data': {'output':
AIMessageChunk(content='Hello! How can I
assist you today?')}}]
```

Chain

Let's revisit the example chain that parsed streaming JSON to explore the streaming events API.

```
chain = model | JsonOutputParser() # This
parser only works with OpenAI right now

events = [
    event
    async for event in chain.astream_events(
        'output a list of the countries
france, spain and japan and their populations
in JSON format. Use a dict with an outer key
of "countries" which contains a list of
countries. Each country should have the key
`name` and `population`',
    version="v1",
    )
]
```

If you examine at the first few events, you'll notice that there are **3** different start events rather than **2** start events.

The three start events correspond to:

- 1. The chain (model + parser)
- 2. The model
- 3. The parser

events[:3]

```
[{'event': 'on_chain_start',
  'run_id': 'aa992fb9-d79f-46f3-a857-
ae4acad841c4',
  'name': 'RunnableSequence',
  'tags': [],
  'metadata': {},
  'data': {'input': 'output a list of the
countries france, spain and japan and their
populations in JSON format. Use a dict with
an outer key of "countries" which contains a
list of countries. Each country should have
the key `name` and `population`'}},
 {'event': 'on_chat_model_start',
  'name': 'ChatOpenAI',
  'run id': 'c5406de5-0880-4829-ae26-
bb565b404e27',
  'tags': ['seq:step:1'],
  'metadata': {},
  'data': {'input': {'messages':
[[HumanMessage(content='output a list of the
countries france, spain and japan and their
populations in JSON format. Use a dict with
an outer key of "countries" which contains a
list of countries. Each country should have
the key `name` and `population`')]]}}},
 {'event': 'on_parser_start',
  'name': 'JsonOutputParser',
```

```
'run_id': '32b47794-8fb6-4ef4-8800-
23ed6c3f4519',
  'tags': ['seq:step:2'],
  'metadata': {},
  'data': {}}]
```

What do you think you'd see if you looked at the last 3 events? what about the middle?

Let's use this API to take output the stream events from the model and the parser. We're ignoring start events, end events and events from the chain.

```
flush=True,
)
if kind == "on_parser_stream":
    print(f"Parser chunk: {event['data']
['chunk']}", flush=True)
    num_events += 1
    if num_events > 30:
        # Truncate the output
        print("...")
        break
```

```
Chat model chunk:
                  1 1
Parser chunk: {}
Chat model chunk: '{\n'
Chat model chunk: ' '
Chat model chunk: ' "'
Chat model chunk: 'countries'
Chat model chunk: '":'
Parser chunk: {'countries': []}
Chat model chunk: ' [\n'
Chat model chunk: ' '
Parser chunk: {'countries': [{}]}
Chat model chunk: ' {\n'
Chat model chunk: ' '
Chat model chunk: ' "'
Chat model chunk: 'name'
Chat model chunk: '":'
Parser chunk: {'countries': [{'name': ''}]}
Chat model chunk:
Parser chunk: {'countries': [{'name':
```

```
'France'}]}
Chat model chunk: 'France'
Chat model chunk: '",\n'
Chat model chunk: ' '
Chat model chunk: ' "'
```

Because both the model and the parser support streaming, we see sreaming events from both components in real time! Kind of cool isn't it?

Filtering Events

Because this API produces so many events, it is useful to be able to filter on events.

You can filter by either component name, component tags or component type.

By Name

```
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries" which contains a list of
countries. Each country should have the key
`name` and `population`',
    version="v1",
    include_names=["my_parser"],
):
    print(event)
    max_events += 1
    if max_events > 10:
        # Truncate output
        print("...")
        break
```

```
{'event': 'on_parser_start', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk': {}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': []}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
```

```
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{}]}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{'name': ''}]}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{'name': 'France'}]}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{'name': 'France',
'population': 670}]}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{'name': 'France',
'population': 670600}]}}}
{'event': 'on_parser_stream', 'name':
'my_parser', 'run_id': '450011c0-6f3b-4ec8-
92d4-6603d9d1d603', 'tags': ['seq:step:2'],
'metadata': {}, 'data': {'chunk':
{'countries': [{'name': 'France',
'population': 67060000}]}}}
```

By Type

```
):
```

```
print(event)
max_events += 1
if max_events > 10:
    # Truncate output
    print("...")
    break
```

```
{'event': 'on_chat_model_start', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'input':
{'messages': [[HumanMessage(content='output a
list of the countries france, spain and japan
and their populations in JSON format. Use a
dict with an outer key of "countries" which
contains a list of countries. Each country
should have the key `name` and
`population`')]]}}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content='')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content='{\n')}}
{'event': 'on_chat_model_stream', 'name':
```

```
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' ')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' "')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content='countries')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content='":')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' [\n')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' ')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
```

```
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' {\n')}}
{'event': 'on_chat_model_stream', 'name':
'model', 'run_id': '9ba1ef9f-5954-4649-b3da-
1171b6abb000', 'tags': ['seq:step:1'],
'metadata': {}, 'data': {'chunk':
AIMessageChunk(content=' ')}}
...
```

By Tags



CAUTION

Tags are inherited by child components of a given runnable.

If you're using tags to filter, make sure that this is what you want.

```
chain = (model |
JsonOutputParser()).with_config({"tags":
    ["my_chain"]})

max_events = 0
async for event in chain.astream_events(
        'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
```

```
"countries" which contains a list of
countries. Each country should have the key
`name` and `population`',
    version="v1",
    include_tags=["my_chain"],
):
    print(event)
    max_events += 1
    if max_events > 10:
        # Truncate output
        print("...")
        break
```

```
{'event': 'on_chain_start', 'run_id':
'd4c78db8-be20-4fa0-87d6-cb317822967a',
'name': 'RunnableSequence', 'tags':
['my_chain'], 'metadata': {}, 'data':
{'input': 'output a list of the countries
france, spain and japan and their populations
in JSON format. Use a dict with an outer key
of "countries" which contains a list of
countries. Each country should have the key
`name` and `population`'}}
{'event': 'on_chat_model_start', 'name':
'ChatOpenAI', 'run_id': '15e46d9f-ccf5-4da2-
b9e3-b2a85873ba4c', 'tags': ['seq:step:1',
'my_chain'], 'metadata': {}, 'data':
{'input': {'messages':
[[HumanMessage(content='output a list of the
countries france, spain and japan and their
```

```
populations in JSON format. Use a dict with
an outer key of "countries" which contains a
list of countries. Each country should have
the key `name` and `population`')]]}}}
{'event': 'on_parser_start', 'name':
'JsonOutputParser', 'run_id': '91945f4f-0deb-
4999-acf0-f6d191c89b34', 'tags':
['seq:step:2', 'my_chain'], 'metadata': {},
'data': {}}
{'event': 'on_chat_model_stream', 'name':
'ChatOpenAI', 'run_id': '15e46d9f-ccf5-4da2-
b9e3-b2a85873ba4c', 'tags': ['seq:step:1',
'my_chain'], 'metadata': {}, 'data':
{'chunk': AIMessageChunk(content='')}}
{'event': 'on_parser_stream', 'name':
'JsonOutputParser', 'run_id': '91945f4f-0deb-
4999-acf0-f6d191c89b34', 'tags':
['seq:step:2', 'my_chain'], 'metadata': {},
'data': {'chunk': {}}}
{'event': 'on_chain_stream', 'run_id':
'd4c78db8-be20-4fa0-87d6-cb317822967a',
'tags': ['my_chain'], 'metadata': {}, 'name':
'RunnableSequence', 'data': {'chunk': {}}}
{'event': 'on_chat_model_stream', 'name':
'ChatOpenAI', 'run_id': '15e46d9f-ccf5-4da2-
b9e3-b2a85873ba4c', 'tags': ['seq:step:1',
'my_chain'], 'metadata': {}, 'data':
{'chunk': AIMessageChunk(content='{"')}}
{'event': 'on_chat_model_stream', 'name':
'ChatOpenAI', 'run_id': '15e46d9f-ccf5-4da2-
b9e3-b2a85873ba4c', 'tags': ['seq:step:1',
```

```
'my_chain'], 'metadata': {}, 'data':
{'chunk':
AIMessageChunk(content='countries')}}
{'event': 'on_chat_model_stream', 'name':
'ChatOpenAI', 'run_id': '15e46d9f-ccf5-4da2-
b9e3-b2a85873ba4c', 'tags': ['seq:step:1',
'my_chain'], 'metadata': {}, 'data':
{'chunk': AIMessageChunk(content='":')}}
{'event': 'on_parser_stream', 'name':
'JsonOutputParser', 'run_id': '91945f4f-0deb-
4999-acf0-f6d191c89b34', 'tags':
['seq:step:2', 'my_chain'], 'metadata': {},
'data': {'chunk': {'countries': []}}}
{'event': 'on_chain_stream', 'run_id':
'd4c78db8-be20-4fa0-87d6-cb317822967a',
'tags': ['my_chain'], 'metadata': {}, 'name':
'RunnableSequence', 'data': {'chunk':
{'countries': []}}}
```

Non-streaming components

Remember how some components don't stream well because they don't operate on **input streams**?

While such components can break streaming of the final output when using astream, astream_events will still yield streaming events from intermediate steps that support streaming!

```
# Function that does not support streaming.
# It operates on the finalizes inputs rather
than
# operating on the input stream.
def extract country names(inputs):
    """A function that does not operates on
input streams and breaks streaming."""
    if not isinstance(inputs, dict):
        return ""
    if "countries" not in inputs:
        return ""
    countries = inputs["countries"]
    if not isinstance(countries, list):
        return
    country_names = [
        country.get("name") for country in
countries if isinstance(country, dict)
    return country_names
chain = (
    model | JsonOutputParser() |
_extract_country_names
) # This parser only works with OpenAI right
now
```

As expected, the astream API doesn't work correctly because _extract_country_names doesn't operate on streams.

```
async for chunk in chain.astream(
    'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries" which contains a list of
countries. Each country should have the key
`name` and `population`',
):
    print(chunk, flush=True)
```

Now, let's confirm that with astream_events we're still seeing streaming output from the model and the parser.

```
num_events = 0

async for event in chain.astream_events(
    'output a list of the countries france,
spain and japan and their populations in JSON
format. Use a dict with an outer key of
"countries" which contains a list of
countries. Each country should have the key
`name` and `population`',
```

```
version="v1",
):
    kind = event["event"]
    if kind == "on chat model stream":
        print(
            f"Chat model chunk:
{repr(event['data']['chunk'].content)}",
            flush=True,
    if kind == "on_parser_stream":
        print(f"Parser chunk: {event['data']
['chunk']}", flush=True)
    num events += 1
    if num events > 30:
        # Truncate the output
        print("...")
        break
```

```
Chat model chunk: ''

Parser chunk: {}

Chat model chunk: '{"'}

Chat model chunk: 'countries'

Chat model chunk: '":'

Parser chunk: {'countries': []}

Chat model chunk: ' [\n'

Chat model chunk: ' '

Parser chunk: {'countries': [{}]}

Chat model chunk: ' {"'}

Chat model chunk: ' {"''}

Chat model chunk: 'name'

Chat model chunk: '":'
```

```
Parser chunk: {'countries': [{'name': ''}]}
Chat model chunk: ' "'
Parser chunk: {'countries': [{'name':
'France'}]}
Chat model chunk: 'France'
Chat model chunk: '",'
                   1 11 1
Chat model chunk:
Chat model chunk: 'population'
Chat model chunk: '":'
Parser chunk: {'countries': [{'name':
'France', 'population': ''}]}
Chat model chunk: ' "'
Parser chunk: {'countries': [{'name':
'France', 'population': '67'}]}
Chat model chunk: '67'
Parser chunk: {'countries': [{'name':
'France', 'population': '67 million'}]}
Chat model chunk: ' million'
Chat model chunk: '"},\n'
```

Propagating Callbacks



A CAUTION

If you're using invoking runnables inside your tools, you need to propagate callbacks to the runnable; otherwise, no stream events will be generated.

(i) NOTE

When using RunnableLambdas or @chain decorator, callbacks are propagated automatically behind the scenes.

```
from langchain_core.runnables import
RunnableLambda
from langchain_core.tools import tool
def reverse_word(word: str):
    return word[::-1]
reverse_word = RunnableLambda(reverse_word)
@tool
def bad_tool(word: str):
    """Custom tool that doesn't propagate
callbacks."""
    return reverse_word.invoke(word)
async for event in
bad_tool.astream_events("hello",
version="v1"):
    print(event)
```

```
{'event': 'on_tool_start', 'run_id':
'39e4a7eb-c13d-46f0-99e7-75c2fa4aa6a6',
'name': 'bad_tool', 'tags': [], 'metadata':
{}, 'data': {'input': 'hello'}}
{'event': 'on_tool_stream', 'run_id':
'39e4a7eb-c13d-46f0-99e7-75c2fa4aa6a6',
'tags': [], 'metadata': {}, 'name':
'bad_tool', 'data': {'chunk': 'olleh'}}
{'event': 'on_tool_end', 'name': 'bad_tool',
'run_id': '39e4a7eb-c13d-46f0-99e7-
75c2fa4aa6a6', 'tags': [], 'metadata': {},
'data': {'output': 'olleh'}}
```

Here's a re-implementation that does propagate callbacks correctly. You'll notice that now we're getting events from the reverse_word runnable as well.

```
@tool
def correct_tool(word: str, callbacks):
    """A tool that correctly propagates
callbacks."""
    return reverse_word.invoke(word,
{"callbacks": callbacks})

async for event in
correct_tool.astream_events("hello",
```

```
version="v1"):
    print(event)
```

```
{'event': 'on_tool_start', 'run_id':
'4263aca5-f221-4eb7-b07e-60a89fb76c5c',
'name': 'correct_tool', 'tags': [],
'metadata': {}, 'data': {'input': 'hello'}}
{'event': 'on_chain_start', 'name':
'reverse_word', 'run_id': '65e3679b-e238-
47ce-a875-ee74480e696e', 'tags': [],
'metadata': {}, 'data': {'input': 'hello'}}
{'event': 'on_chain_end', 'name':
'reverse_word', 'run_id': '65e3679b-e238-
47ce-a875-ee74480e696e', 'tags': [],
'metadata': {}, 'data': {'input': 'hello',
'output': 'olleh'}}
{'event': 'on_tool_stream', 'run_id':
'4263aca5-f221-4eb7-b07e-60a89fb76c5c',
'tags': [], 'metadata': {}, 'name':
'correct_tool', 'data': {'chunk': 'olleh'}}
{'event': 'on_tool_end', 'name':
'correct_tool', 'run_id': '4263aca5-f221-
4eb7-b07e-60a89fb76c5c', 'tags': [],
'metadata': {}, 'data': {'output': 'olleh'}}
```

If you're invoking runnables from within Runnable Lambdas or @chains, then callbacks will be passed automatically on your behalf.

```
from langchain_core.runnables import
RunnableLambda
```

```
async def reverse_and_double(word: str):
    return await reverse_word.ainvoke(word) *
2

reverse_and_double =
RunnableLambda(reverse_and_double)

await reverse_and_double.ainvoke("1234")

async for event in
reverse_and_double.astream_events("1234",
version="v1"):
    print(event)
```

```
{'event': 'on_chain_start', 'run_id':
'714d22d4-a3c3-45fc-b2f1-913aa7f0fc22',
'name': 'reverse_and_double', 'tags': [],
'metadata': {}, 'data': {'input': '1234'}}
{'event': 'on_chain_start', 'name':
'reverse_word', 'run_id': '35a6470c-db65-
4fe1-8dff-4e3418601d2f', 'tags': [],
'metadata': {}, 'data': {'input': '1234'}}
{'event': 'on_chain_end', 'name':
'reverse_word', 'run_id': '35a6470c-db65-
```

```
4fe1-8dff-4e3418601d2f', 'tags': [],
'metadata': {}, 'data': {'input': '1234',
'output': '4321'}}
{'event': 'on_chain_stream', 'run_id':
'714d22d4-a3c3-45fc-b2f1-913aa7f0fc22',
'tags': [], 'metadata': {}, 'name':
'reverse_and_double', 'data': {'chunk':
'43214321'}}
{'event': 'on_chain_end', 'name':
'reverse_and_double', 'run_id': '714d22d4-a3c3-45fc-b2f1-913aa7f0fc22', 'tags': [],
'metadata': {}, 'data': {'output':
'43214321'}}
```

And with the @chain decorator:

```
from langchain_core.runnables import chain

@chain
async def reverse_and_double(word: str):
    return await reverse_word.ainvoke(word) *
2

await reverse_and_double.ainvoke("1234")

async for event in
reverse_and_double.astream_events("1234",
```

```
version="v1"):
    print(event)
```

```
{'event': 'on_chain_start', 'run_id':
'17c89289-9c71-406d-90de-86f76b5e798b',
'name': 'reverse_and_double', 'tags': [],
'metadata': {}, 'data': {'input': '1234'}}
{'event': 'on_chain_start', 'name':
'reverse_word', 'run_id': 'b1105188-9196-
43c1-9603-4f2f58e51de4', 'tags': [],
'metadata': {}, 'data': {'input': '1234'}}
{'event': 'on_chain_end', 'name':
'reverse_word', 'run_id': 'b1105188-9196-
43c1-9603-4f2f58e51de4', 'tags': [],
'metadata': {}, 'data': {'input': '1234',
'output': '4321'}}
{'event': 'on_chain_stream', 'run_id':
'17c89289-9c71-406d-90de-86f76b5e798b',
'tags': [], 'metadata': {}, 'name':
'reverse_and_double', 'data': {'chunk':
'43214321'}}
{'event': 'on_chain_end', 'name':
'reverse_and_double', 'run_id': '17c89289-
9c71-406d-90de-86f76b5e798b', 'tags': [],
'metadata': {}, 'data': {'output':
'43214321'}}
```